

# **Supervised Learning for NLP I**

HSS 510: NLP for HSS

Taegyeon Kim

Mar 27, 2024

# Agenda

## Things to be covered

- Overview of supervised machine learning
- Step 1: Building a labeled data set
- Step 2: Extracting features
- Step 3: Selecting and training model(s)
- Step 4: Evaluating performance
- Guided labeling: text classification with movie review data in Python

# Supervised Learning

We will focus on (text) classification with supervised learning

- Goal
  - To classify documents into pre-defined categories
  - E.g., sentiment of comments, stance on policy issues, topic of news articles, etc.
- We need
  - Human-labeled data set
  - Model (algorithm) that maps documents (i.e., their features) to labels
  - Evaluation approaches
    - Performance metrics, cross-validation, etc.

# Supervised Learning

## Supervised vs. Unsupervised

	Supervised	Unsupervised
Objective	Trained on a labeled data to learn a mapping from input to output	Find patterns or structures within data without labels
Outcome	Pre-defined categories	Not quite pre-defined
Model evaluation	Explicit metrics such as accuracy, precision, recall, or MSE	Can involve qualitative assessment
Examples	Classification/regression for texts	Topic models

# Supervised Learning

## Regression vs. Classification

- Regression
  - The outcome of interest is continuous or ordered (beyond binary)
  - E.g., OLS regression (+ non-linear regression algorithms such as random forest regression)
- Classification
  - The outcome is a value in an unordered set (i.e. categories)
- The two approaches share the broad principles of supervised learning and can be adapted

# Supervised Learning vs. Dictionary Methods

## Limitations of dictionary methods

- Lack of learning (as in the the name machine “learning”)
  - (Largely) ignores context
    - Polysemy, co-occurrences/interactions, etc.
    - Interactions are effectively modeled in random forest, deep neural networks, and large language models
- Therefore, (generally) suboptimal performance

# Overview of Process

## Broad process

- Step 1: build a labeled data set
- Step 2: extract features
- Step 3: select and train model(s)
- Step 4: evaluate performance

# Overview of Process

## Step 1: build a labeled data set

- Documents with human-annotated labels (a.k.a. ground-truth) :  $C$
- Randomly split into a training set and a test set:  $C = C_{train} + C_{test}$
- E.g., identifying YouTube comments containing hate speech
  - $C$ : 10,000 comments labeled for the presence of hate speech
  - $C_{training}$  : 8,000 comments for training
  - $C_{test}$  : 2,000 comments for test



# Overview of Process

Step 1: build a labeled data set

Doc number	Text	y
1	This is great!	0
2	%@% ***k off!	1
...		
9999	This is sick	0
10000	Love BTS <3	0

# Overview of Process

## Step 2: extract features

- Generate  $X_{train}$  (feature matrix) from  $C_{train}$  (train set)
- E.g., count vectors, TF-IDF, or embeddings

Index	Token 1	Token 2	...	Token V-1	Token V
1	3	1.4	...	1.7	6
2	-0.8	6.4	...	5.7	-1.6
...					
7999	-2.8	0.9	...	3.3	-0.6
8000	3.7	1.4	...	5.7	-5.8

# Overview of Process

Step 3: select/train model(s)

Index	y	Token 1	Token 2	...	Token V-1	Token V
1	0	3	1.4	...	1.7	6
2	1	-0.8	6.4	...	5.7	-1.6
...						
7999	0	-2.8	0.9	...	3.3	-0.6
8000	0	3.7	1.4	...	5.7	-5.8

# Overview of Process

## Step 3: select/train model(s)

- Choose a model  $F$  (e.g., logistic regression) and learn model parameters  $\beta$  (e.g., an array of coefficients)
  - The model provides a mapping between  $X_{train}$  and  $y_{train}$
- Loss (cost) function: measures how much model predictions ( $\hat{y}_{train}$ ) differ from the true labels ( $y_{train}$ )
  - $\beta$  is estimated in a way that minimizes the difference
  - $\hat{y}_{train} = F(\hat{\beta} * X_{train})$
- As a result, we get a classifier:  $\hat{y} = F(\hat{\beta} * X)$

# Overview of Process

Step 3: select/train model(s)

Index	$y$	$\hat{y}$	Token 1	Token 2	...	Token V-1	Token V
1	0	0	3	1.4	...	1.7	6
2	1	1	-0.8	6.4	...	5.7	-1.6
...							
7999	0	0	-2.8	0.9	...	3.3	-0.6
8000	0	0	3.7	1.4	...	5.7	-5.8

# Overview of Process

## Step 4: evaluate performance

- We held out another labeled set  $C_{test}$  ( $n = 2,000$ ) (**why?**)
- Use the classifier  $F(\hat{\beta} * X)$  to generate predictions  $\hat{y}_{test}$
- Compare the predictions  $\hat{y}_{test}$  and the true labels  $y_{test}$
- Performance metrics include accuracy, precision, recall, etc.
- *(Then use the classifier for unlabeled data)*

# Overview of Process

Step 4: evaluate performance

Index	$y$	$\hat{y}$	Token 1	Token 2	...	Token V-1	Token V
1	1	1	3.12	1.99	...	5.77	0.36
2	1	0	-0.8	1.14	...	9.71	-1.66
...							
1999	0	0	-2.11	0.95	...	1.23	-0.62
2000	0	0	3.71	1.48	...	1.7	-5.84

# Bias, Variance, and Overfitting

## Bias




- The degree to which the model's predictions deviate from the true labels in a systematic manner
- A model with high bias make predictions that are consistently off-target

## Variance


- The degree to which the model generalizes to different data
- High variance means low generalizability




# Bias, Variance, and Overfitting




# Bias, Variance, and Overfitting



**overfitting**



**underfitting**




**Good balance**


# Bias, Variance, and Overfitting

## Overfitting and underfitting


- If a model learns the training data “too well” (low bias), it can lead to overfit
- This happens when the model mistakes noise for signal
- The model would not generalize to the test set (high variance)



**overfitting**



**underfitting**



**Good balance**

# Bias, Variance, and Overfitting

## Training-test split

- A minimal measure to prevent overfitting
- The primary goal here is to make our model as generalizable as possible
- “Generalizable” means being able to perform well on unseen documents (other than the documents the model was trained on)
- When a model learns the noise or random fluctuations in the training set, this typically results in a model that performs poorly on new the test set

# Prediction and Explanation

How are they different?

- Predictive modeling emphasizes predictive performance and generalizability (i.e., out-of-sample prediction)
- Explanatory modeling emphasizes hypothesis the statistical significance of an individual coefficient(s) associated with a (relational) hypothesis
- See [Shmueli \(2010\)](#) for a detailed treatment of the differences between prediction and explanation

# Step 1: Building a Labeled Data Set

How do we obtain a labeled data set?

- A form of manual content analysis
- “Ground truth” or “gold standard” fed to machines
- Our decisions/labels reflect latent features linked to the categories (some of which we are unconscious of)
- Manually labeled data are used for training (train set) and evaluation (test set)

# Step 1: Building a Labeled Data Set

How do we obtain a labeled data set?

- Expert labeling
  - In many projects, a few domain experts work on a labeling (after training)
  - Annotators are trained to learn the concept and related guidelines
  - E.g., a researcher + two RAs from the department
- Crowd-sourced labeling
  - “Wisdom of crowd”: aggregated judgments of (online) non-experts converge to judgments of experts at much lower cost ([Benoit et al, 2016](#))
  - Difficult to educate annotators on sophisticated tasks
  - Inductive measurement based on loose conceptualization

# Step 1: Building a Labeled Data Set

## Expert labeling vs. Crowd Sourcing

- Deductive vs. inductive
- Degree of training
- Scalability (cost)



# Step 1: Building a Labeled Data Set

## Selected texts for manual labeling

- Should reflect the entire corpus
- Mismatch leading to low performance: shift/drift
- E.g., drift in anti-vaccine discourse throughout 2020

# Step 1: Building a Labeled Data Set

## Iterative process

- Definition/operationalization does not often take place at once but in an iterative process
- In many cases, it is difficult to specify an entire annotation guidelines ex ante
- Preliminary labeling rule are written and applied to an initial set of docs
  - Annotators identify ambiguities in the rule
  - Revision of the rule → ...

# Step 1: Building a Labeled Data Set

## Dealing with subjectivity

- Many concepts in humanities and social sciences are not straightforward
- They can involve high levels of subjectivity
- This is, from the beginning, why 1) careful conceptualization and 2) writing an excellent labeling rule, and 3) training coders are extremely important
- **Evaluation metrics:** Krippendorff's  $\alpha$ , Cohen's  $\kappa$  (alternatives include Pearson's  $r$ , Spearman's  $\rho$ ) (recommended R package: [irr](#))

# Step 1: Building a Labeled Data Set

## Who are the annotators?

- Expert coding
  - Academics/students ([Javdani and Chang 2023](#))
- Crowdsourcing
  - Skewed distribution of worked hours ([Difallah et al. 2018](#))
  - Inattentive workers ([Peyton et al. 2022](#); [Ternovski 2022](#))
  - LLM-based responses ([Veselovsky et al. 2023](#))
  - Demographic characteristic ([Al Kuwatly et al. 2020](#))

## Step 2: Extract Features

$$C_{train} \rightarrow X_{train}$$

- Options include count vectors, TF-IDF vectors, word/document embeddings, etc.

Index	Token 1	Token 2	...	Token V-1	Token V
1	3	1.4	...	1.7	6
2	-0.8	6.4	...	5.7	-1.6
...					
7999	-2.8	0.9	...	3.3	-0.6
8000	3.7	1.4	...	5.7	-5.8

## Step 3: Select/train Model(s)

So far we have:

- Built a labeled data set (Step 1)
- Generated a feature matrix (Step 2)
- This means that we have the outcome ( $y$ ) and features ( $X_{train}$ )

Now we will:

- Select a model  $F$
- Learn model parameters  $\beta$  to build a classifier ( $\hat{y} = F(\hat{\beta} * X)$ )

## Step 3: Select/train Model(s)

Numerous algorithms

- Logistic regression
- Naive Bayes
- Support vector machine
- Tree-based models (decision tree, random forest, XGBoost, etc.)
- Neural networks
- Etc.

## Step 3: Select/train Model(s)

### Logistic regression

- Used to classify a document into binary categories
  - Multinomial logistic regression for more than two
- One of the most useful analytics tools in science (not just NLP/ML)
- The baseline supervised learning algorithm for classification
- Forms the basis of neural networks



## Step 3: Select/train Model(s)

Components of logistic regression ( $j$  documents  $n$  features)

- Features (e.g., tokens)
  - A document is represented as a vector of features  $\vec{x} = [x_1, \dots, x_n]$
- A classification function ( $F$ )
  - $p(y = 1 | x)$  is computed for each document given the feature vector and  $\beta$
- Loss function and algorithm for optimizing it (gradient descent)

## Step 3: Select/train Model(s)

How does logistic regression compute predicted probabilities?

- $p(y = 1 | x)$ 
  - We want to know the probability of  $y = 1$  given a feature vector  $x \vec{=} [x_1, \dots, x_n]$
  - For a simple count vector, it would be the number times each token appears in the document
- Logistic regression learns  $\beta$ , a vector of coefficients
  - A bias term  $b$ : a single number (a.k.a. intercept)
  - Weights  $w \vec{=} [w_1, \dots, w_n]$ 
    - E.g., features signaling hateful intention would get high weights
  - With  $b$ ,  $w \vec{}$ , and  $x \vec{}$ , we compute  $z (= (\sum_{i=1}^n w_i x_i) + b)$


## Step 3: Select/train Model(s)

How does logistic regression compute predicted probabilities?

- $z = (\sum_{i=1}^n w_i x_i) + b = \vec{w} \cdot \vec{v} + b$

- $\sigma(z) = \frac{1}{1+e^{-z}} = \frac{1}{1+\exp(-z)}$

### Step 3: Select/train Model(s)



## Step 3: Select/train Model(s)

How does logistic regression compute predicted probabilities?

$$p(y = 1 | x) = \sigma(\vec{w} \cdot \vec{x} + b)$$

$$p(y = 0 | x) = 1 - \sigma(\vec{w} \cdot \vec{x} + b)$$

## Step 3: Select/train Model(s)

How do predicted probabilities turn into binary labels ( $\hat{y}$ )?

$$\begin{cases} 1 & \text{if } P(y = 1 | x) > 0.5 \\ 0 & \text{otherwise} \end{cases}$$

## Step 3: Select/train Model(s)

Sentiment classification from movie reviews

*“It’s hokey. There are virtually no surprises, and the writing is second-rate. So why was it so enjoyable? For one thing, the cast is great, Another nice touch is the music. I was overcome with the urge to get off the couch and start dancing? It sucked me in, and it’ll do the same to you.”*

## Step 3: Select/train Model(s)

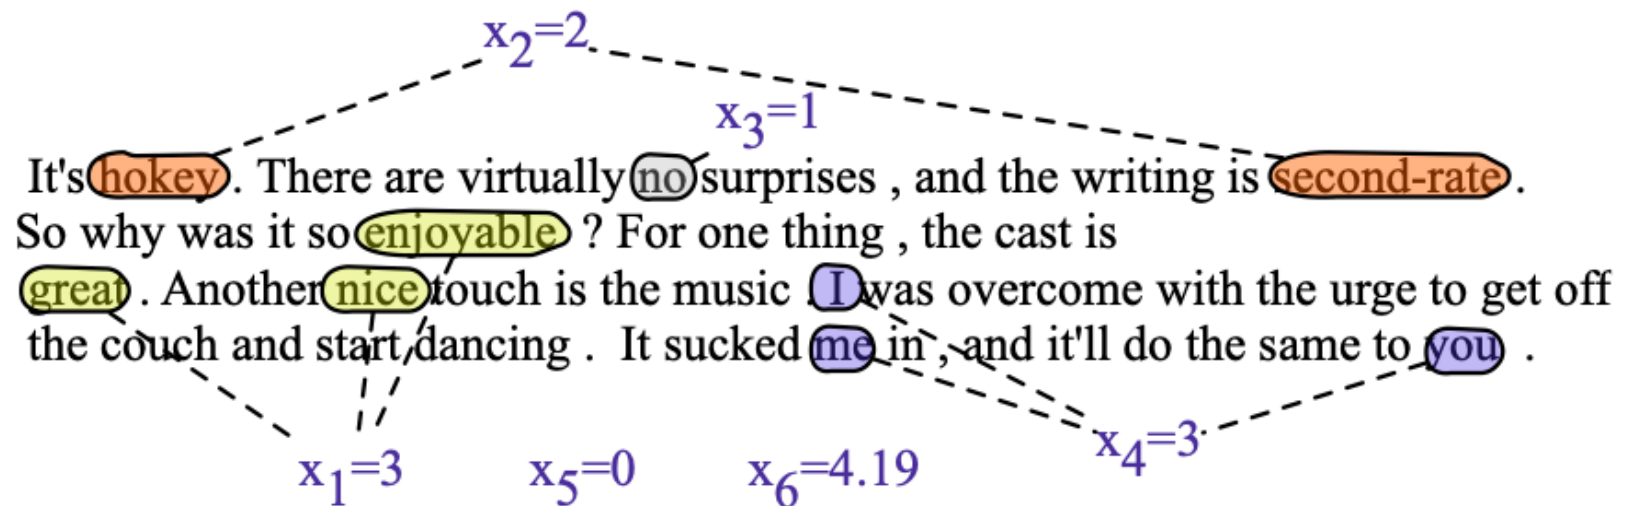
Sentiment classification from movie reviews

Var	Definition
$x_1$	count(positive lexicon) $\in$ doc)
$x_2$	count(negative lexicon) $\in$ doc)
$x_3$	$\begin{cases} 1 & \text{if "no"} \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$
$x_4$	count(1st and 2nd pronouns $\in$ doc)
$x_5$	$\begin{cases} 1 & \text{if "!"} \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$
$x_6$	log(word count of doc)



## Step 3: Select/train Model(s)

Sentiment classification from movie reviews



## Step 3: Select/train Model(s)

### Positive

- $$\begin{aligned} p(+x) &= p(y = 1 | x) = \sigma(\vec{w} \cdot \vec{x} + b) \\ &= \sigma([2.5, -5.0, -1.2, 0.5, 2.0, 0.7] \cdot [3, 2, 1, 3, 0, 4.19] + 0.1) \\ &= \sigma(.833) \\ &= 0.70 \end{aligned}$$

### Negative

- $$\begin{aligned} p(-x) &= p(Y = 0 | x) = 1 - \sigma(\vec{w} \cdot \vec{x} + b) \\ &= 0.30 \end{aligned}$$

## Step 3: Select/train Model(s)

### Learning coefficients

- MLE (Maximum Likelihood Estimation)
- Loss function / cross entropy
- Gradient descent

## Step 4: Evaluate Performance

We have

- Manually labeled documents
- Split them into  $C_{train}$  (training set) and  $C_{test}$  (test set)
- Trained a classifier on  $C_{train}$  (with  $y_{train}$  and  $X_{train}$ )  $\rightarrow F(\hat{\beta}^* X)$

Now we need to evaluate its performance on  $C_{test}$

- We compare  $\hat{y}_{test}$  (predicted labels) against  $y_{test}$  (true labels)

# Step 4: Evaluate Performance

## Performance metrics

- Accuracy: the proportion of all predictions (both positive and negative) that the model got right
- Precision: the proportion of positive predictions that were actually correct
- Recall: the proportion of actual positives that were correctly predicted
- F-1: the harmonic (as opposed to arithmetic) mean of precision and recall

## Step 4: Evaluate Performance

Confusion matrix: predictions against true labels

		True condition	
		Positive	Negative
Prediction	Positive	True Positive	False Positive (Type I error)
	Negative	False Negative (Type II error)	True Negative

## Step 4: Evaluate Performance

$$\text{Accuracy: } \frac{TP+TN}{TP+TN+FP+FN}$$

		True condition	
		Positive	Negative
Prediction	Positive	True Positive	False Positive (Type I error)
	Negative	False Negative (Type II error)	True Negative

## Step 4: Evaluate Performance

Precision:  $\frac{TP}{TP+FP}$

		True condition	
		Positive	Negative
Prediction	Positive	True Positive	False Positive (Type I error)
	Negative	False Negative (Type II error)	True Negative



## Step 4: Evaluate Performance


Recall:  $\frac{TP}{TP+FN}$

		True condition	
		Positive	Negative
Prediction	Positive	True Positive	False Positive (Type I error)
	Negative	False Negative (Type II error)	True Negative

## Step 4: Evaluate Performance

$$F-1: (2 \times \textit{precision} \times \textit{recall}) / (\textit{precision} + \textit{recall})$$

- Why not arithmetic mean  $((\textit{precision} + \textit{recall})/2)$ ?



# Step 4: Evaluate Performance

Precision/recall/F-1 & accuracy

- 100 positives
- 80 predicted positives
- 60 true positives

		True condition		
		Positive	Negative	
Prediction	Positive	60		80
	Negative			
		100		

## Step 4: Evaluate Performance

Precision/recall/F-1 & accuracy

- Precision:  $\frac{60}{60+20} = 0.75$

		True condition		
		Positive	Negative	
Prediction	Positive	60	20	80
	Negative			
		100		

## Step 4: Evaluate Performance

Precision/recall/F-1 & accuracy

- Recall:  $\frac{60}{60+40} = 0.6$

		True condition		
		Positive	Negative	
Prediction	Positive	60	20	80
	Negative	40		
		100		

# Step 4: Evaluate Performance

## Precision/recall/F-1 & accuracy

- Precision:  $\frac{60}{60+20} = 0.75$
- Recall:  $\frac{60}{60+40} = 0.6$
- Accuracy:  $\frac{60+50}{60+20+40+50} = 0.65$

		True condition		
		Positive	Negative	
Prediction	Positive	60	20	80
	Negative	40	50	
		100		

# Step 4: Evaluate Performance

## Precision/recall/F-1 & accuracy

- Precision:  $\frac{60}{60+20} = 0.75$
- Recall:  $\frac{60}{60+40} = 0.6$
- Accuracy:  $\frac{60+150}{60+20+40+150} = 0.78$

		True condition		
		Positive	Negative	
Prediction	Positive	60	20	80
	Negative	40	150	
		100		

# Step 4: Evaluate Performance

An extremely imbalanced case

- Accuracy: ??
- Precision: ??
- Recall: ??
- F-1: ??

		True condition		
		Positive	Negative	
Prediction	Positive	2	1	3
	Negative	8	989	997
		10	990	100



# Step 4: Evaluate Performance

An extremely imbalanced case

- Accuracy: 0.991
- Precision: 0.66
- Recall: 0.2
- F-1: 0.31

		True condition		
		Positive	Negative	
Prediction	Positive	2	1	3
	Negative	8	989	997
		10	990	100

# Step 4: Evaluate Performance

## Reminders

- Random train-test split:  $C_{train}$  ,  $C_{test}$ 
  - E.g., 10,000 comments labeled for hate speech into 8,000 and 1,000
- Our classifier learned **parameters**, maximizing performance on  $C_{train}$  and evaluating it on  $C_{test}$


# Step 4: Evaluate Performance

## Parameters vs. hyper-parameters


- Parameter
  - Learned (estimated) from data (internal to the model)
  - E.g., logistic regression weights/coefficients ( $= \beta$ )
- Hyper-parameters
  - Defines the model structure itself (not internal to the model)
  - E.g., the size of a regularization term in logistic regression, the number of layers or learning rate in neural networks, etc.

## Step 4: Evaluate Performance

Shallow neural network



Deep neural network



## Step 4: Evaluate Performance

Hyper-parameters influence model performance, and we want to “tune” them

- With different hyper-parameter values, we could fit a model configured with each value on  $C_{train}$  and evaluate performance on  $C_{test}$
- E.g., regularization strength  $\lambda$  in logistic regression
  - Train different models with different  $\lambda$  values on  $C_{train}$
  - Evaluate on  $C_{test}$
  - Pick the best performing model

## Step 4: Evaluate Performance

What could go wrong?

- In comparing different models (different hyperparameter values), we might overfit on  $C_{train}$
- By repeatedly using the training set, our comparison can be affected by the specific characteristics of the test set

# Step 4: Evaluate Performance

## Validation set

- We split the labeled data set into a training set, a “validation set”, and a test set:  $C_{train}$  ,  $C_{validation}$  , and  $C_{test}$
- Train a model on  $C_{train}$  and see how it performs on  $C_{validation}$
- Repeat this step for multiple configurations (e.g., hyper-parameters)
- Pick the best-performing configuration
- Train a final model based on the configuration on  $C_{train} + C_{validation}$
- Evaluate on  $C_{test}$

# Step 4: Evaluate Performance

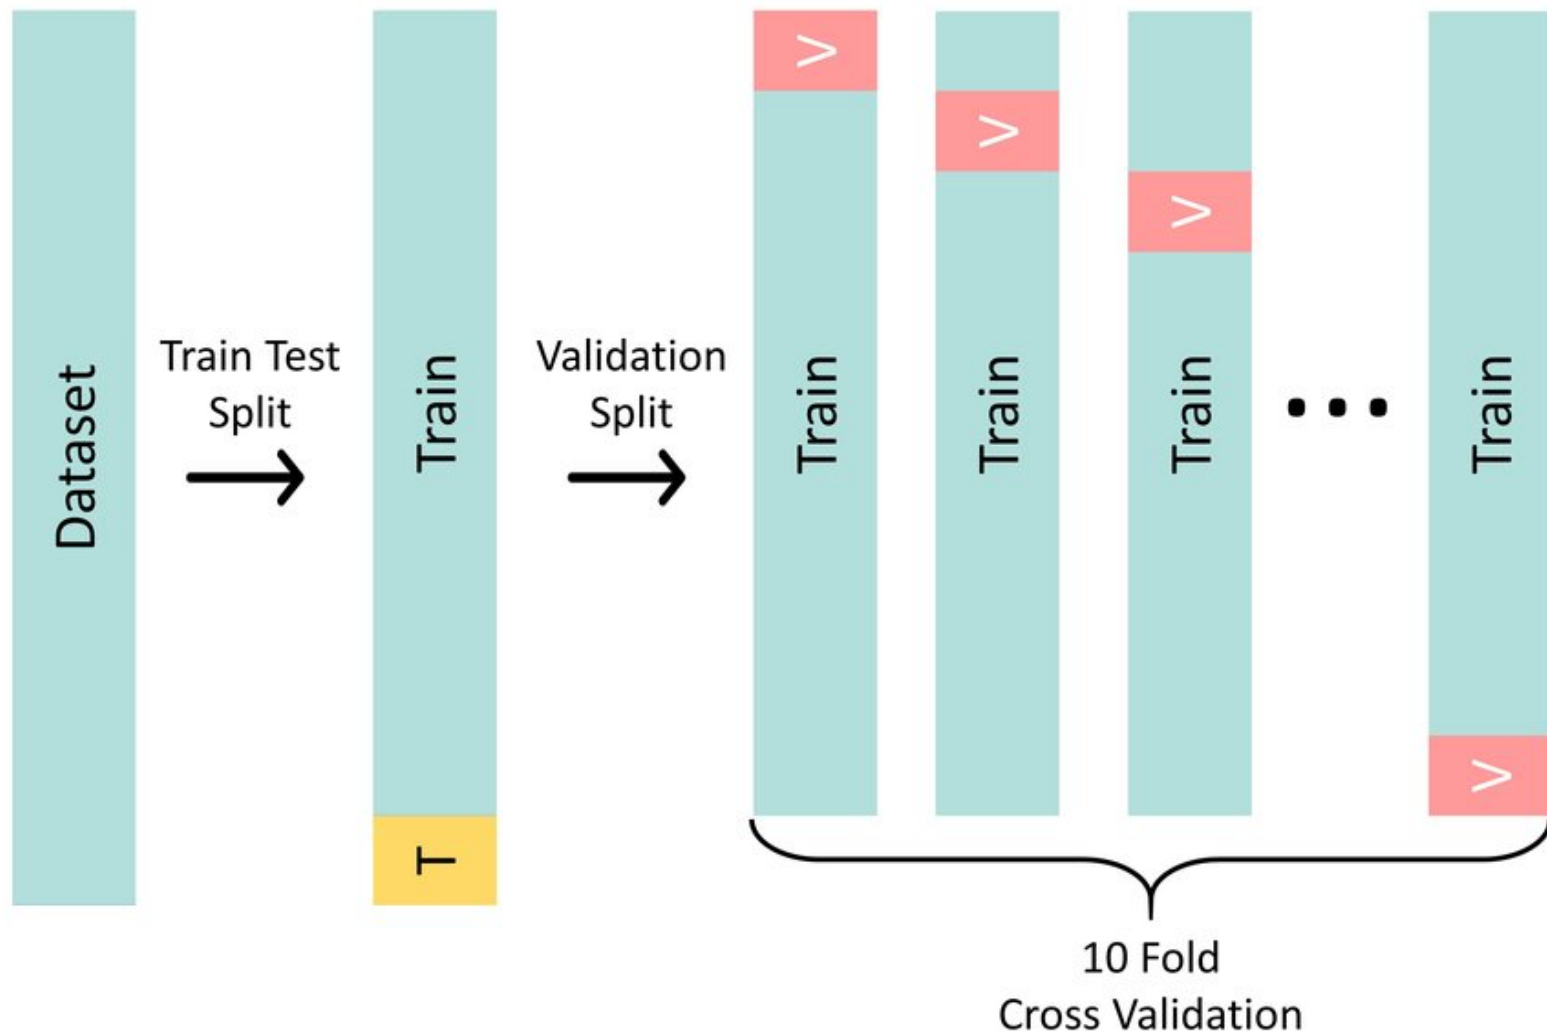
## $K$ -fold cross-validation

- Randomly split  $C_{train}$  into  $K$  equal parts or “folds” (commonly 5 or 10)
- For each iteration
  - Treat one fold as the “validation set”
  - Train your model on the remaining  $K - 1$  folds
  - Evaluate performance on the validation set kept aside
- After cycling through all iterations
  - Aggregate the performance metrics obtained from each iteration
  - Choose the classifier with the highest cross-validated performance
  - This step may involve not just hyperparameter tuning but also things like feature representation, etc.
- (Re)train the chosen best classifier on  $C_{train}$  (all  $K$  folds combined) and evaluate on  $C_{test}$





## Step 4: Evaluate Performance



# Summary

Supervised text classification provides a highly useful tool to assign labels to documents

- Be aware of the principles of building a labeled data set
  - Conceptualization, intercoder reliability, annotator bias, etc.
- Validate, validate, and validate!
- Choose appropriate evaluation metrics

# Guided Coding

Text classification with movie reviews data ([link](#))